

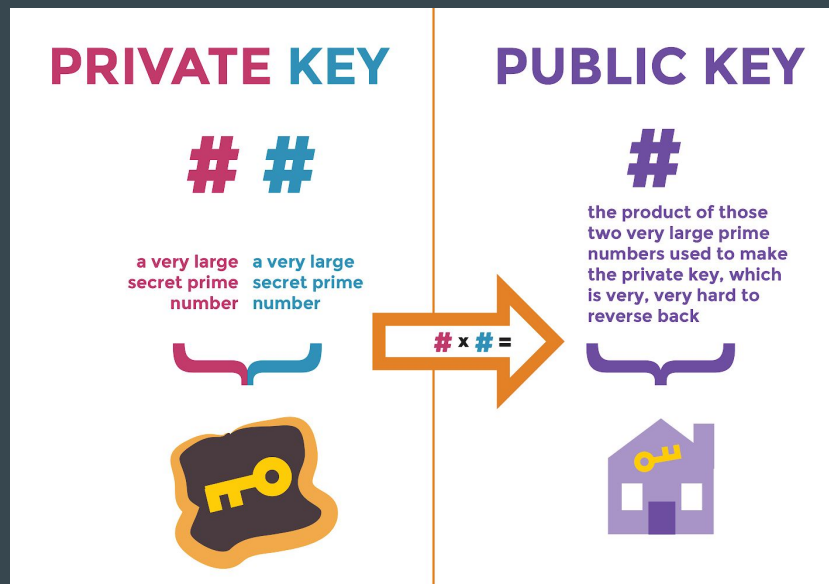
Widening the Sieve

...

Richard Muniu, Mickey Haregot, Kevin Zheng

Finding Prime Numbers

- Prime numbers hard to find efficiently
- This makes them useful in security protocols
 - E.g. public/private key encryption
 - Rivest-Shamir-Aldeman (RSA) algorithm
- 2000 year-old algorithms still surprisingly effective today



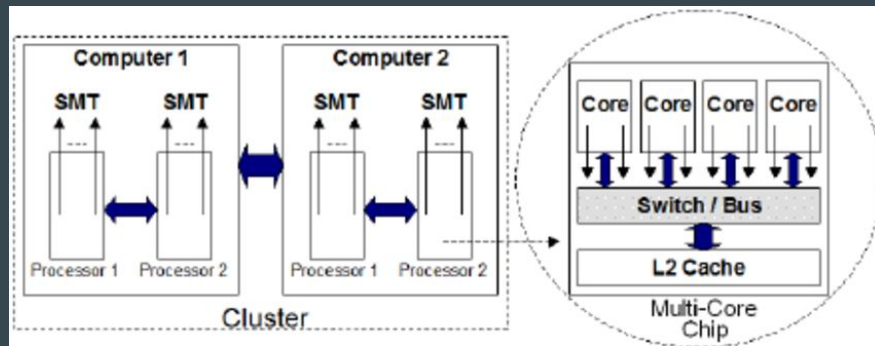
The Sieve of Eratosthenes

1. Have an array for each integer between 0 and some max N
2. Iterate through the array, whenever you find an “unmarked” element, mark all of its multiples.
3. All unmarked elements by end of the run are prime numbers
4. Features: unbalanced load, not embarrassingly parallel, low computation per step

	2	3	4	5	6	7	8	9	10	Prime numbers
11	12	13	14	15	16	17	18	19	20	
21	22	23	24	25	26	27	28	29	30	
31	32	33	34	35	36	37	38	39	40	
41	42	43	44	45	46	47	48	49	50	
51	52	53	54	55	56	57	58	59	60	
61	62	63	64	65	66	67	68	69	70	
71	72	73	74	75	76	77	78	79	80	
81	82	83	84	85	86	87	88	89	90	
91	92	93	94	95	96	97	98	99	100	
101	102	103	104	105	106	107	108	109	110	
111	112	113	114	115	116	117	118	119	120	


System Traits Desired

- Decided to use single machine
 - No communication
 - Used single node on cluster (Strelka)
- Reasons why
 - Sieve algorithm relatively simple
 - Using multiple nodes result in large communication costs
 - Results required to be aggregated at end of runs
 - May become large part of total runtime



Methods of Implementation

- Sequential
 - Baseline test
- OpenMP
 - Parallelized by base: each thread investigates a “unmarked” number separately
- CUDA
 - Massively parallel version of OpenMP implementation.
 - Takes advantage of SIMD architecture
- Balanced Pthreads
 - Each thread is responsible for a different section of the array and the base process broadcasts “bases” for the other threads to investigate



OpenMP

- Uses OpenMP's loop primitives
- Marks multiples in parallel
- Load balancing problem: heavy work assigned to one thread
- Resolved using dynamic scheduling of iteration chunks

	2	3	4	5	6	7	8	9	10	Prime numbers
11	12	13	14	15	16	17	18	19	20	
21	22	23	24	25	26	27	28	29	30	
31	32	33	34	35	36	37	38	39	40	
41	42	43	44	45	46	47	48	49	50	
51	52	53	54	55	56	57	58	59	60	
61	62	63	64	65	66	67	68	69	70	
71	72	73	74	75	76	77	78	79	80	
81	82	83	84	85	86	87	88	89	90	
91	92	93	94	95	96	97	98	99	100	
101	102	103	104	105	106	107	108	109	110	
111	112	113	114	115	116	117	118	119	120	

```
schedule(static):
```

```
*****
```

```
*****
```

```
*****
```

```
*****
```

```
schedule(dynamic):
```

```
*  ** ** * * * *      * *  **  * * * *      * * *
```

```
*      *      *  **  * *  *      * *  * *
```


```
*      *      *  **  * *  *      * *  * * * *
```

```
* *      *  **  * *  *      * *  * *  * *
```

CUDA

- Similar to OpenMP implementation
- Massively parallel
- Every possible “base” has its multiples simultaneously marked off
- Takes advantage of SIMT architecture
- Multiple kernel calls to combat max-thread limitations

```
if (threadIdx.x < 4) {  
    A;  
    B;  
} else {  
    X;  
    Y;  
}  
Z;
```



Balanced pthreads

- Splits up array into segments
- Master-worker model
- Meant to balance workload
- Master discovers “bases” and broadcasts them to workers
- Broadcast uses shared variable and barriers

	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	100

Testing

- Compared sequential vs parallel runtimes on Strelka
- OpenMP/balanced pthreads: tested both seq_sieve and par_sieve
 - 32-core CPU
 - hi-mem partition
- CUDA: ran cudaSieve on
 - 4x NVIDIA 2080 Ti GPUs
 - gpu-02 partition.
- For both OpenMP and balanced pthread, dynamic allocation in RAM failed for arrays of size 10^{13} and above.
- For CUDA, dynamic allocation in DRAM failed for arrays of size 10^{12} and above.
- Used average values from three runs of seq_sieve and par_sieve for each method.



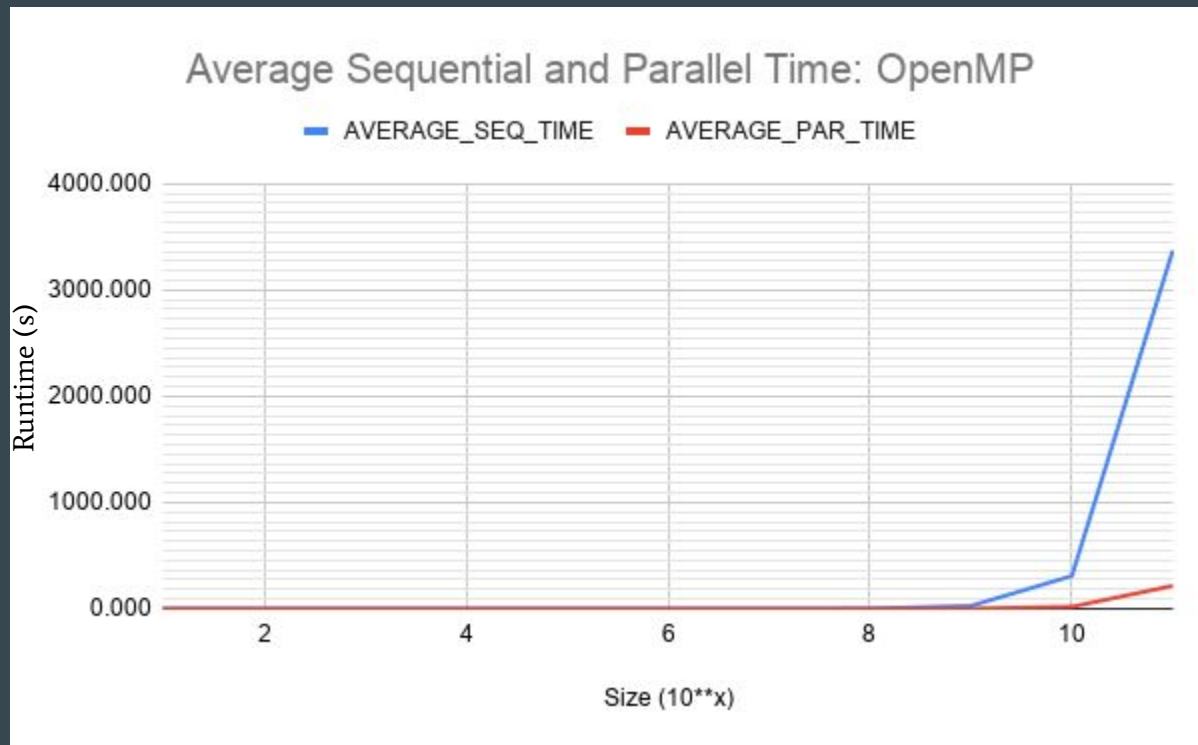
RESULTS AND DISCUSSION

OpenMP

- Nonlinear speedup of up to 16
- No speedup until arrays of size 10^7

Size (10^x)	Average Sequential Time(s)	Average Parallel Time(s)	Speedup
1	0.001	0.017	0.058
2	0.001	0.016	0.063
3	0.001	0.017	0.059
4	0.001	0.016	0.061
5	0.001	0.018	0.057
6	0.008	0.017	0.462
7	0.073	0.027	2.750
8	2.050	0.201	10.182
9	27.495	1.923	14.298
10	308.771	19.372	15.939
11	3371.608	216.177	15.596

OpenMP

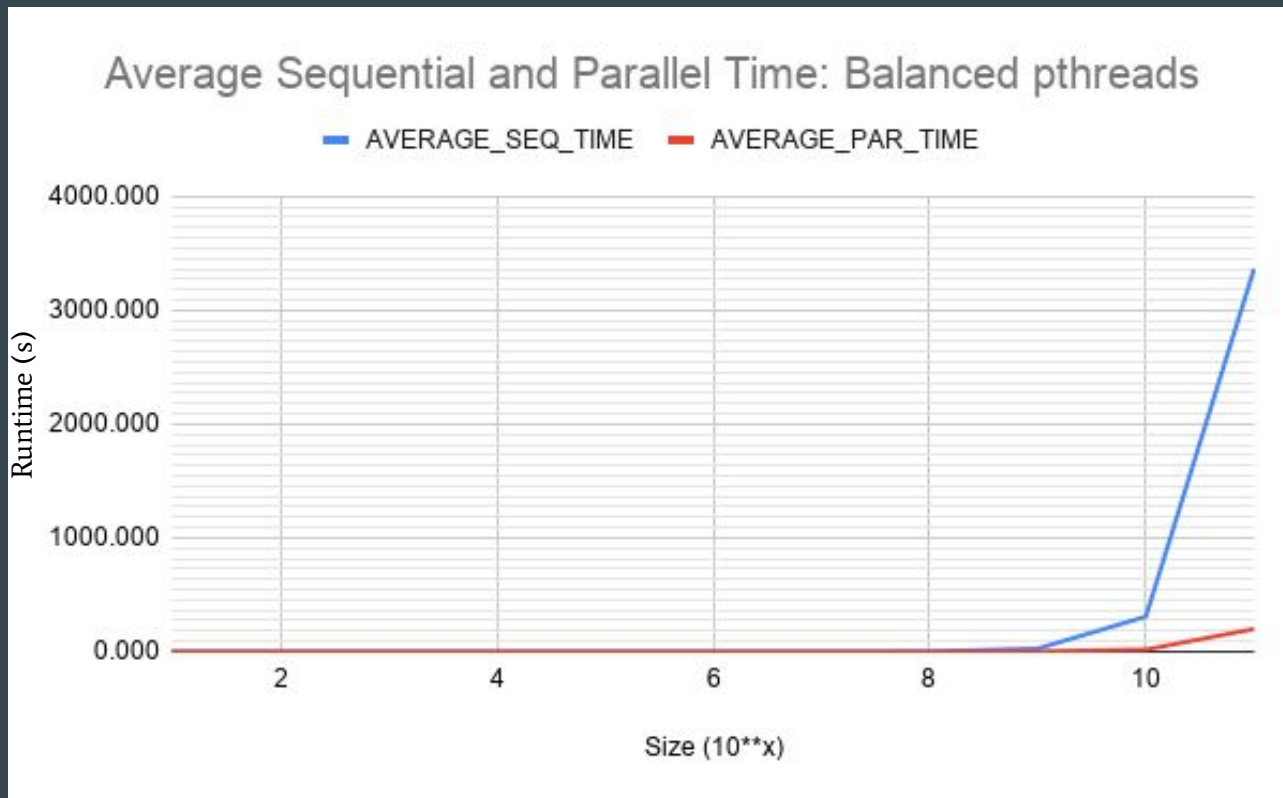


Balanced pthreads

- Nonlinear speedup of close to 17
- No consistent speedup until arrays of size 10^7

Size (10^x)	Average Sequential Runtime	Average Parallel Runtime	Speedup
1	0.005	0.002	2.500
2	0.001	0.002	0.429
3	0.001	0.003	0.375
4	0.001	0.005	0.200
5	0.001	0.009	0.154
6	0.007	0.016	0.458
7	0.071	0.046	1.547
8	2.204	0.228	9.653
9	27.518	2.440	11.278
10	308.638	18.786	16.429
11	3364.497	199.944	16.827

Balanced pthreads

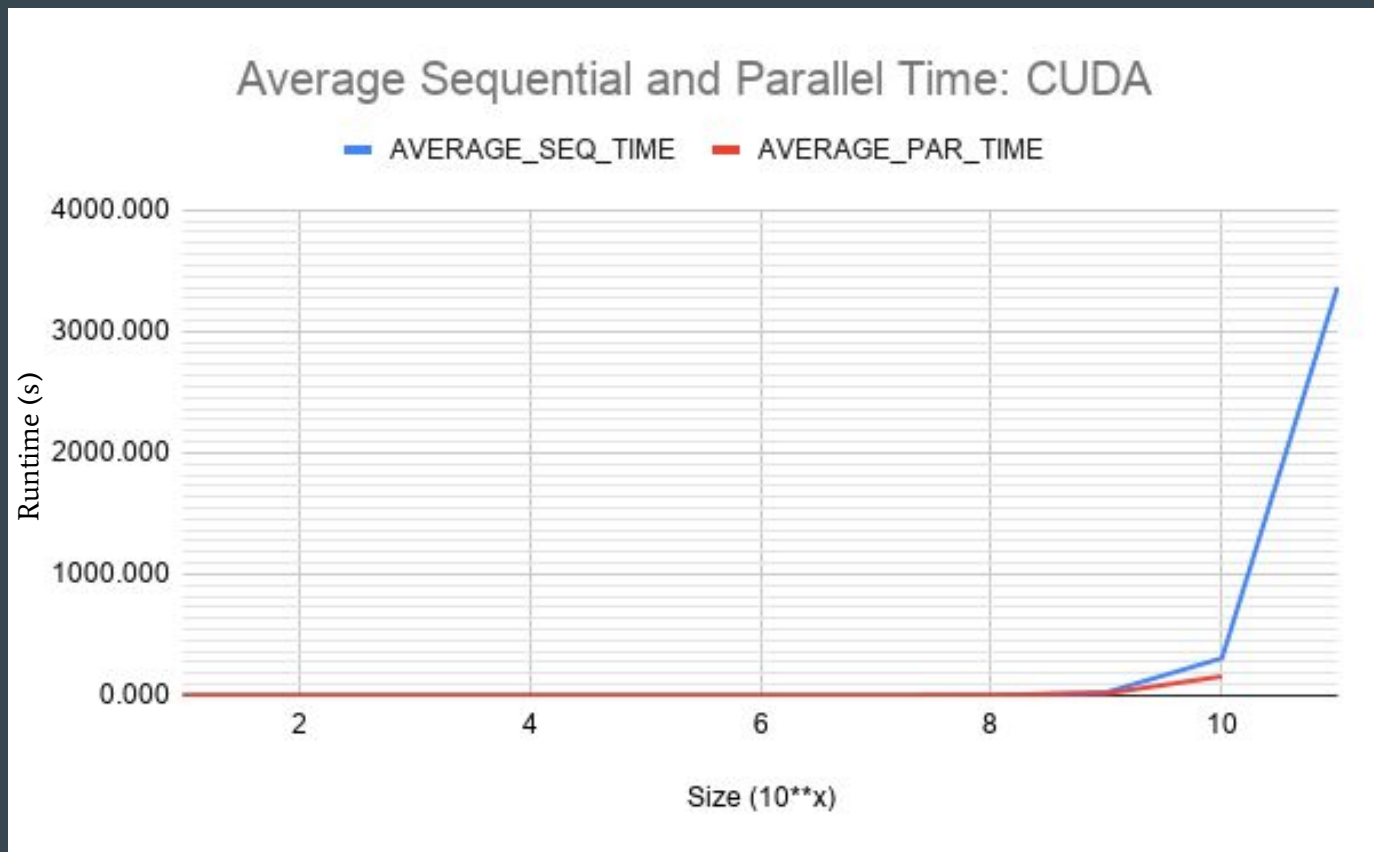


CUDA

- Nonlinear speedup of up to 2
- No consistent speedup until arrays of size 10^9 , but relatively consistent runtime until that point
- Lower performance than both other implementations

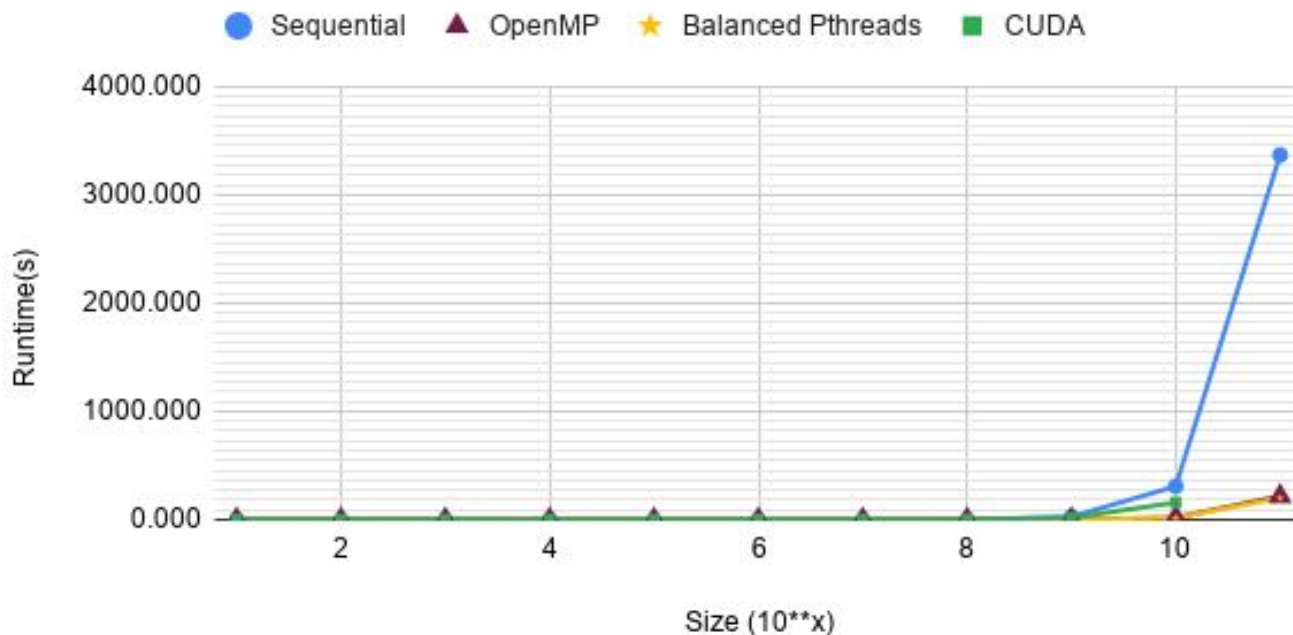
Size (10^x)	AVERAGE SEQ TIME	AVERAGE PAR TIME	Speedup
1	0.005	3.128	0.002
2	0.001	3.039	0.000
3	0.001	3.060	0.000
4	0.001	3.057	0.000
5	0.001	3.076	0.000
6	0.007	3.090	0.002
7	0.071	3.295	0.021
8	2.204	4.536	0.486
9	27.518	17.802	1.546
10	308.638	158.962	1.942
11	-	-	-

CUDA



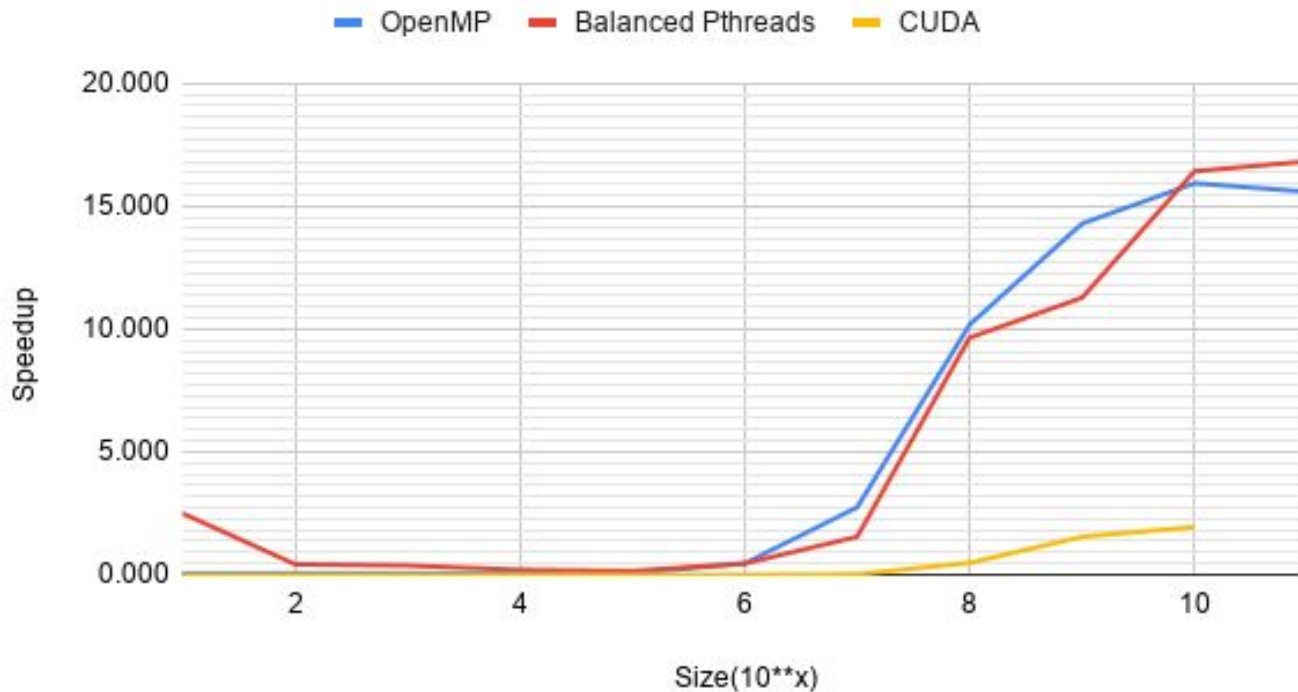
Runtime Trend and Comparison

Sequential, OpenMP, Balanced Pthreads and CUDA Runtime Comparison



Speedup Trend and Comparison

OpenMP, Balanced Pthreads and CUDA Speedup Comparison



Conclusion and Future Directions

- A lot to gain from parallelization in form of speedup and efficiency
- Benefits are not obvious/present at low task sizes due to parallelization overheads
- CUDA performed slower than other methods
 - hardware limitations?
 - Potential software optimizations?
- Some future directions to consider:
 - Further optimizations to CUDA implementation
 - Explore distributed memory approach to eliminate max array size limitation



QUESTIONS

kzhengl@swarthmore.edu

mharegol@swarthmore.edu

rmuniul@swarthmore.edu